Viswesh Kesarla

Hyderabad, India | +91 7995292102 | visweshkesarla@gmail.com
linkedin.com/in/visweshkesarla | github.com/visweshdev | visweshkesarla.vercel.app

Education

Woxsen University, Hyderabad, India 2024 – 2028 (Expected)
Bachelor of Technology (B.Tech) in Computer Science and Engineering; GPA: 8.61/10
Relevant Coursework: Data Structures and Algorithms, Object-Oriented Programming, Database Management Systems, Computer Networks, AI/ML Foundations

Experience

AI/ML Intern, Wipro, Hyderabad, India May 2026 – July 2026

- Built a Retrieval-Augmented Generation (RAG) pipeline in Python with LangChain, Hugging Face sentence embeddings, and ChromaDB, indexing 1,000+ chunks from 25+ PDFs for semantic question answering.
- Prototyped LLM workflows with Google ADK to orchestrate document retrieval and question-answering steps.
- Debugged retrieval and generation failures through iterative testing of chunking strategies, Top-K settings, embeddings, and system prompts, improving answer relevance and reducing hallucinations.

Projects

ReMem: Outcome-Learning Memory for AI Agents | GitHub 2026
Python, Model Context Protocol (MCP), FAISS, ChromaDB, LangChain, NumPy

- Built an open-source memory layer for AI agents that learns from task outcomes which memories help, retires outdated facts, and fits context to a token budget via knapsack selection.
- Shipped as a Python API, MCP server (Claude Code/Desktop), and LangChain retriever with FAISS and ChromaDB backends; 85 tests running in GitHub Actions CI.
- Benchmarked against similarity search over 15 held-out seeds with ablations: +13.8 percentage points task success (+10.2 with real embeddings) using 20% fewer context tokens.

PaperForge AI 2026
Python, RAG, Large Language Models (LLMs), FAISS

- Built an end-to-end RAG pipeline with semantic retrieval and FAISS vector search to extract methodologies from academic research papers.
- Applied semantic and recursive character splitting to preserve context boundaries, improving Recall@K on dense documents.
- Built an LLM-based research-to-code parser using strict system prompts and output schemas to turn unstructured formulas into compilable Python functions.

URL Shortener | GitHub 2026
Python, Django, PostgreSQL, Docker, GitHub Actions

- Built a full-stack Django web application to generate, store, and redirect custom short links, with link previews and click analytics.
- Containerized with Docker Compose (PostgreSQL, WhiteNoise, multi-environment settings) and added a CI test stage running flake8 and code coverage in an isolated container.

Skills

Languages: Python, SQL, TypeScript
AI/ML: Machine Learning, Natural Language Processing (NLP), Large Language Models (LLMs), RAG, LangChain, LangGraph, Google ADK, Model Context Protocol (MCP), Hugging Face Transformers, sentence-transformers, Ollama, Prompt Engineering, LLM Evaluation, PyTorch, Scikit-learn, NumPy
Databases and Backend: ChromaDB, FAISS, PostgreSQL, MongoDB, Django, FastAPI
Tools: Git, GitHub Actions, CI/CD, Docker, Docker Compose, pytest, Jupyter Notebook, Google Colab

Certifications

- Anthropic: Generative AI and Agentic Developer Suite (Agent Skills, Model Context Protocol, Claude API/Code), 2026
- Git Fundamentals, LearnkartS (2026); Data Structures Using Python, Packt (2025); Programming for Everybody (Python), University of Michigan (2024)

Hackathons

- Udgham 24-Hour Hackathon, Woxsen University (2024, 2025)